

In [1]: `!pip install -q kaggle`

In [2]: `from google.colab import files`
`files.upload()`

Choose Files No file chosen

Upload widget is only available when the

cell has been executed in the current browser session. Please rerun this cell to enable.

Saving kaggle.json to kaggle.json

Out[2]: `{'kaggle.json': b'{"username":"dhananjaythakur001","key":"bc71203698055d3957e4f5270361eba4"}'}`

In [3]: `!mkdir -p ~/.kaggle`
`!cp kaggle.json ~/.kaggle/`
`!chmod 600 ~/.kaggle/kaggle.json`

In [4]: `!kaggle datasets download -d salader/dogs-vs-cats`

Dataset URL: <https://www.kaggle.com/datasets/salader/dogs-vs-cats>

License(s): unknown

In []: `!unzip dogs-vs-cats.zip`

Archive: dogs-vs-cats.zip

replace dogs_vs_cats/test/cats/cat.10.jpg? [y]es, [n]o, [A]ll, [N]one, [r]ename:

In [6]: `import tensorflow as tf`
`from tensorflow import keras`
`from tensorflow.keras.models import Sequential`
`from tensorflow.keras.layers import Dense, Conv2D, Flatten, Dropout, MaxPooling2D,`
`from tensorflow.keras.preprocessing.image import ImageDataGenerator`

In [7]: `train_ds = keras.utils.image_dataset_from_directory(`
 `directory = '/content/train',`
 `labels = 'inferred',`
 `label_mode = 'binary', # Changed to 'binary'`
 `batch_size = 32,`
 `image_size = (256, 256)`
`)`
`test_ds = keras.utils.image_dataset_from_directory(`
 `directory = '/content/test',`
 `labels = 'inferred',`
 `label_mode = 'binary', # Changed to 'binary'`
 `batch_size = 32,`
 `image_size = (256, 256)`
`)`

Found 20000 files belonging to 2 classes.

Found 5000 files belonging to 2 classes.

In [8]: `from tensorflow.keras.utils import to_categorical`
`def process(image, label):`
 `image = tf.cast(image/255. , tf.float32)`

```

label = to_categorical(label, num_classes=2) # One-hot encode the labels
return image, label

train_ds = train_ds.map(process)
test_ds = test_ds.map(process)

```

```

In [9]: model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(256,256,3))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2), strides = 2, padding = 'valid'))
model.add(Conv2D(64, kernel_size=(3, 3), activation='relu', padding = 'valid'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2), strides = 2, padding = 'valid'))
model.add(Conv2D(128, kernel_size=(3, 3), activation='relu', padding = 'valid'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2), strides = 2, padding = 'valid'))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(2, activation='sigmoid'))
model.summary()

```

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:
107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first layer in
the model instead.

```

```

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

Model: "sequential"

Layer (type)	Output Shape	
conv2d (Conv2D)	(None, 254, 254, 32)	
batch_normalization (BatchNormalization)	(None, 254, 254, 32)	
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	
conv2d_1 (Conv2D)	(None, 125, 125, 64)	
batch_normalization_1 (BatchNormalization)	(None, 125, 125, 64)	
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 64)	
conv2d_2 (Conv2D)	(None, 60, 60, 128)	
batch_normalization_2 (BatchNormalization)	(None, 60, 60, 128)	
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 128)	
flatten (Flatten)	(None, 115200)	
dense (Dense)	(None, 128)	1
dropout (Dropout)	(None, 128)	
dense_1 (Dense)	(None, 64)	
dropout_1 (Dropout)	(None, 64)	
dense_2 (Dense)	(None, 2)	



Total params: 14,848,258 (56.64 MB)

Trainable params: 14,847,810 (56.64 MB)

Non-trainable params: 448 (1.75 KB)

```
In [10]: model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
history = model.fit(train_ds, epochs=10, validation_data=test_ds)
```

Epoch 1/10

625/625 ————— **67s** 85ms/step - accuracy: 0.5612 - loss: 1.5967 - val_accuracy: 0.6026 - val_loss: 0.6627

Epoch 2/10

625/625 ————— **70s** 82ms/step - accuracy: 0.6432 - loss: 0.6416 - val_accuracy: 0.6334 - val_loss: 0.6301

Epoch 3/10

625/625 ————— **51s** 82ms/step - accuracy: 0.7120 - loss: 0.5632 - val_accuracy: 0.7486 - val_loss: 0.5139

Epoch 4/10

625/625 ————— **82s** 83ms/step - accuracy: 0.7763 - loss: 0.4782 - val_accuracy: 0.7822 - val_loss: 0.4532

Epoch 5/10

625/625 ————— **88s** 93ms/step - accuracy: 0.8154 - loss: 0.4112 - val_accuracy: 0.7928 - val_loss: 0.4472

Epoch 6/10

625/625 ————— **77s** 86ms/step - accuracy: 0.8524 - loss: 0.3485 - val_accuracy: 0.7824 - val_loss: 0.4932

Epoch 7/10

625/625 ————— **68s** 109ms/step - accuracy: 0.8864 - loss: 0.2742 - val_accuracy: 0.7112 - val_loss: 0.7170

Epoch 8/10

625/625 ————— **53s** 85ms/step - accuracy: 0.9206 - loss: 0.2075 - val_accuracy: 0.7620 - val_loss: 0.5551


Epoch 9/10

625/625 ————— **55s** 88ms/step - accuracy: 0.9525 - loss: 0.1295 - val_accuracy: 0.7932 - val_loss: 0.7786

Epoch 10/10

625/625 ————— **52s** 83ms/step - accuracy: 0.9690 - loss: 0.0876 - val_accuracy: 0.7866 - val_loss: 1.3534

```
In [11]: # Plotting Graph - accuracy
import matplotlib.pyplot as plt
plt.plot(history.history['accuracy'], label='train acc')
plt.plot(history.history['val_accuracy'], label='test acc')
plt.legend()
plt.show()
```



```
In [12]: # Plot Confusion Matrix
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report
import numpy as np

# Extract true Labels from the test dataset
y_true = []
for images, labels in test_ds:
    y_true.extend(np.argmax(labels.numpy(), axis=1)) # Convert to class indices

Y_pred = model.predict(test_ds)
y_pred = np.argmax(Y_pred, axis=1)
print('Confusion Matrix')

# Use y_true instead of test_ds.classes
print(confusion_matrix(y_true, y_pred))

print('Classification Report')
target_names = ['Cats', 'Dogs']
print(classification_report(y_true, y_pred, target_names=target_names))
```

157/157 ————— 9s 54ms/step

Confusion Matrix

[[1532 968]
[1553 947]]

Classification Report

	precision	recall	f1-score	support
Cats	0.50	0.61	0.55	2500
Dogs	0.49	0.38	0.43	2500
accuracy			0.50	5000
macro avg	0.50	0.50	0.49	5000
weighted avg	0.50	0.50	0.49	5000


```
In [14]: import cv2
def predict_image(image_path):
    img = cv2.imread(image_path)
    img = cv2.resize(img, (256, 256))
    img = img / 255.0
    img = np.expand_dims(img, axis=0)
    prediction = model.predict(img)
    predicted_class = np.argmax(prediction)
    if predicted_class == 0:
        print("The image is predicted as a cat.")
    else:
        print("The image is predicted as a dog.")

# Replace 'path/to/your/image.jpg' with the actual path to your image
predict_image('/content/dogs_vs_cats/test/cats/cat.10.jpg')
```

1/1 ————— 2s 2s/step

The image is predicted as a cat.

```
In [26]: test_img = cv2.imread('/content/dogs_vs_cats/test/cats/cat.10030.jpg')
plt.imshow(test_img)
plt.show()
test_img.shape
test_img = cv2.resize(test_img, (256, 256))
test_input = test_img.reshape((1, 256, 256, 3))
model.predict(test_input)
```



1/1 ————— 0s 49ms/step

Out[26]: array([[1., 0.]], dtype=float32)

```
In [27]: def preprocess_image(image_path):
            image = cv2.imread(image_path) # Load image
            image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB) # Convert BGR to RGB
            image = cv2.resize(image, (256, 256)) # Resize to match model input
            image = image / 255.0 # Normalize pixel values
            image = np.expand_dims(image, axis=0) # Add batch dimension
            return image

        def predict_image(image_path):
            image = preprocess_image(image_path)
            prediction = model.predict(image)[0][0] # Get prediction score
            label = "Cat" if prediction > 0.5 else "Dog"
            confidence = max(prediction, 1 - prediction) * 100 # Convert to percentage
            print(f"Prediction: {label} (Confidence: {confidence:.2f}%)")
            image_path = "/content/dogs_vs_cats/test/cats/cat.10030.jpg" # Replace with your i
            predict_image(image_path)
```

1/1 ————— 0s 106ms/step

Prediction: Cat (Confidence: 100.00%)

In []: